PERFORMANCE ASSESSMENT

KAILI HAMILTON

Masters of Science in Data Analytics, Western Governors University

Course: D208 Predictive Modeling

Instructor: Dr. Daniel Smith

Program Mentor: Krissy Bryant

October, 2023

TABLE OF CONTENTS

Performance Assessment	1
A1	3
A2	3
B1	3
B2	3
B3	3
C1	4
C2	9
C31	4
C42	1
C52	3
D12	4
D22	4
D32	5
E12	8
E22	8
E3	4
F1	4
F2	7
G3	8
Н	8
I3	8
J	9

A1

My research question is, "What factors predict how much data per year a customer uses on average?"

A2

The goals of the data analysis is to determine which factors influence how much data a customer uses per year on average and how strong the relationship is among the each factor and the response variable, "Bandwidth_GB_Year". This will help the company optimize the services they offer to customers based on these factors and can thus save money. Also, the company may want to offer a data plan for their services with different thresholds of data.

B1

Assumptions of multiple linear regression include (Middleton, nd) (Nair, 2021):

- 1. Linearity between the explanatory and response variables,
- 2. Residuals of the regression are normally distributed,
- 3. Little to no multicollinearity exists in the data,
- 4. Homoscedasticity
- 5. Independence of observations

B2

Two benefits of using Python in support of various phases of the analysis include, but are not limited to, computing power to calculate and fit a regression model using vast amounts of data and creating myriad data visualizations to inform the analysis.

B3

Multiple linear regression is an appropriate technique to answer the research question (summarized in parts A1 and A2) because I am striving to find relationships among variables. Specifically, I want to predict which factors influence my continuous response variable "Bandwidth_GB_Year" (Van den Broeck, nd).

C1

My data cleaning goals are designed to help me answer my research question (Part A1). My data cleaning goals are to treat null values, remove extreme outliers, re-express categorical variables for use in multiple linear regression, and to determine which variables to use in my initial regression model (Middleton, nd).

Treated null values.

There were no null values in the data.

Check for null values

<pre>df.isna().sum()</pre>	
Population	0
Age	0
Income	0
Churn	0
Outage_sec_perweek	0
Yearly_equip_failure	0
Techie	0
Port_modem	0
Tablet	0
InternetService	0
Phone	0
Multiple	0
OnlineSecurity	0
OnlineBackup	0
StreamingTV	0
StreamingMovies	0
Tenure	0
MonthlyCharge	0
Bandwidth_GB_Year	0
dtype: int64	

no null values

Removed outliers.

Regression is sensitive to outliers. However, due to the natural variation of the distribution of some variables, I decided to keep most outliers.

The variables "Population", "Income", "Outage_sec_perweek", and "Yearly_equip_failure" all contain outliers. I created boxplots for all numerical variables, calculated the number of outliers for each, and graphed a scatterplot against the response variable, "Bandwidth_GB_Year" to inform whether to drop outliers. The regression plots helped to determine if the outliers were influencing the overall trend of the data. Also, I did not want to change the shape of the distribution so I opted to only drop the most extreme outliers. See the copy of the code for further details regarding code.

POPULATION

```
: sns.boxplot(x='Population', data=df, showmeans=True)
: <Axes: xlabel='Population'>
                                                                                                AGE
                                                                                                sns.boxplot(x='Age', data=df, showmeans=True)
                                                                                                <Axes: xlabel='Age'>
                20000
                             40000
                                         60000
                                                     80000
                                                                 100000
                                    Population
  def find_outliers_IQR(df):
     q1=df.quantile(0.25)
     q3=df.quantile(0.75)
     IQR=q3-q1
     outliers = df[((df<(q1-1.5*IQR)) | (df>(q3+1.5*IQR)))]
     return outliers
  outliers = find_outliers_IQR(df['Population'])
print('number of outliers: " + str(len(outliers)))
print('max outlier value: ' + str(outliers.max()))
print('min outlier value: ' + str(outliers.min()))
                                                                                                                                                                                     90
                                                                                                      20
                                                                                                                  30
                                                                                                                             40
                                                                                                                                        50
                                                                                                                                                   60
                                                                                                                                                               70
                                                                                                                                                                         80
                                                                                                                                           Age
  number of outliers: 937
max outlier value: 111850
min outlier value: 31816
                                                                                                no outliers in age
          sns.regplot(x='Income', y='Bandwidth_GB_Year', data=df, ci=None)
           <Axes: xlabel='Income', ylabel='Bandwidth_GB_Year'>
                7000
                6000
                5000
            Bandwidth_GB_Year
                4000
                3000
                2000
                1000
                      0
                                           50000
                             0
                                                             100000
                                                                               150000
                                                                                                  200000
                                                                                                                    250000
                                                                        Income
```

Outliers that were dropped:

• Population > 100,000

```
df.drop(df[df['Population'] > 100000].index, inplace=True)
df.shape
(9991, 19)
df['Population'].describe()
         9991.000000
         9730.486037
mean
std
        14341.542493
min
           0.000000
25%
          737.500000
50%
         2905.000000
75%
       13161.000000
max
       98660.000000
Name: Population, dtype: float64
```

Income > 200,000

```
df['Income'].describe()
          9994.000000
count
         39710.411833
       27861.200412
std
min
           348.670000
         19219.835000
25%
50%
         33156.205000
75%
        53226.895000
       196746.000000
max
Name: Income, dtype: float64
df.drop(df[df['Population'] > 100000].index, inplace=True)
df.shape
(9991, 19)
```

• Yearly_equip_failure = 6

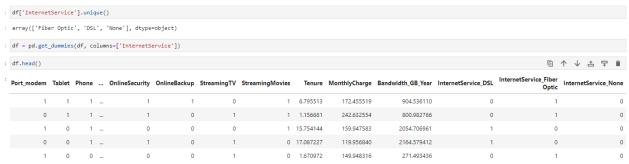
```
df.drop(df[df['Yearly_equip_failure'] == 6].index, inplace=True)
df.shape
(9999, 19)
df['Yearly_equip_failure'].describe()
        9999.000000
count
           0.397440
mean
std
           0.633512
min
           0.000000
25%
           0.000000
50%
           0.000000
75%
           1.000000
           4.000000
max
Name: Yearly_equip_failure, dtype: float64
```

Re-expressed categorical data types.

First I re-expressed <u>ordinal</u> categorical data types using ordinal encoding. These included variables with a "yes" or "no" value. All "yes" values were replaced with "1" and all "no" values were replaced with "0". Then I made sure the datatype changed from object to int64.

```
: df['Churn'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
: df['Churn'].value_counts()
       7341
        2650
                                                                                    df.dtypes
  Name: Churn, dtype: int64
                                                                                     Population
: df['Techie'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
                                                                                                                int64
                                                                                    Income
                                                                                                              float64
  df['Port_modem'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
                                                                                    Churn
                                                                                                                int64
                                                                                    Outage sec perweek
                                                                                                              float64
: df['Tablet'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
                                                                                     Yearly_equip_failure
  df['Phone'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
                                                                                     Techie
                                                                                                                int64
                                                                                    Port modem
                                                                                                                int64
                                                                                    Tablet
                                                                                                                int64
: df['Multiple'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['OnlineSecurity'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
df['OnlineBackup'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
df['StreamingTV'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
                                                                                    InternetService
                                                                                                               object
                                                                                    Multiple
                                                                                                                int64
                                                                                    OnlineSecurity
                                                                                                                int64
  df['StreamingMovies'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
                                                                                                                int64
                                                                                    OnlineBackup
                                                                                    StreamingTV
                                                                                                                int64
                                                                                     StreamingMovies
                                                                                                                int64
: df['StreamingMovies'].value_counts()
                                                                                                              float64
                                                                                    Tenure
                                                                                    MonthlyCharge
                                                                                                              float64
  0
        5104
                                                                                     Bandwidth GB Year
                                                                                                              float64
                                                                                    dtype: object
  Name: StreamingMovies, dtype: int64
```

Next, I re-expressed <u>nominal</u> categorical data types for the variable "InternetService" using pd.getdummies(). "InternetService" has three values, with no inherent rank.



Determined which variables to use in my initial model.

I initially kept all variables that I thought would influence the response variable "Bandwidth_GB_Year". "Job" was one of the variables I thought to include, but there 639 unique values, so I excluded it from the analysis. I also dropped "Martial" for the same reason, too many unique values. I also dropped "Children" so the focus is just the customer, not the household. This is also discussed in Part C4.



The explanatory variables used in my initial regression are:

- Population
- Age
- Income
- Churn
- Outage_sec_perweek
- Yearly_equip_failure
- Techie
- Port_modem
- Tablet
- InternetService
- Phone
- Multiple
- OnlineSecurity
- OnlineBackup
- StreamingTV
- StreamingMovies
- Tenure
- MonthlyCharge

C₂

The dependent (aka response) variable for the data analysis is "Bandwidth_GB_Year". From the summary statistics displayed below we see that there are 9991 values with a range of approximately 7003 GB. On average, customers use about 3391 GB per year with a median value of 3261 GB. The middle 50% of customers use between 1236 and 5585 GB per year.

```
df['Bandwidth GB Year'].describe()
         9991.000000
count
         3390.795380
mean
std
        2185.252241
min
         155.506715
25%
        1236.046551
50%
         3260.745232
75%
         5584.704954
         7158.981530
max
Name: Bandwidth_GB_Year, dtype: float64
```

The independent (aka explanatory) variables in this analysis are listed below along with their summary statistics. There are 9,991 entries for each variable.

• <u>Population</u>. The distribution of Population is highly skewed right as evidenced by the mean population of 9730 being much larger than the median population, 2905. The middle 50% of customers live in a one mile radius population size between 737 and 13,161 people.

```
df['Population'].describe()
          9991.000000
count
          9730.486037
mean
         14341.542493
std
min
             0.000000
25%
           737.500000
50%
          2905.000000
75%
         13161.000000
         98660.000000
max
Name: Population, dtype: float64
```

• Age. The mean and median age for customers is 53, with the oldest customer 89 years old and the youngest 20 years old.

```
df['Age'].describe()
        9991.000000
count
         53.082174
mean
std
          20.702085
          18.000000
min
25%
          35.000000
50%
          53.000000
75%
          71.000000
          89.000000
max
Name: Age, dtype: float64
```

• <u>Income</u>. On average, customers have a yearly income of about \$39,716 (median value \$33,169). The income of the middle 50% of customers ranges between \$19,215 and \$53,228.

```
df['Income'].describe()
count
         9991.000000
mean
         39715.501730
        27863.826217
std
min
           348.670000
25%
        19214.740000
50%
         33168.880000
        53227.795000
75%
        196746.000000
max
Name: Income, dtype: float64
```

• <u>Churn</u>. Out of 9991 customers in the dataset, 73.5% of customers churned within the last month, leaving 26.5% of customers staying with the company.

```
df['Churn'].value_counts()

Churn
0 7341
1 2650

Name: count, dtype: int64
```

• <u>Outage_sec_perweek</u>. On average, customers' neighborhood's experience 10 seconds of outage time per week.

```
df['Outage_sec_perweek'].describe()
count
        9991.000000
         10.002278
mean
std
           2.976494
           0.099747
min
25%
           8.019310
50%
          10.019720
75%
          11.971418
          21.207230
Name: Outage sec perweek, dtype: float64
```

 Yearly equip failure. The average number of times per year a customer's equipment failed and replaced was 0.4, with a median of 0 times. The maximum yearly equipment failure a customer experienced was 4.

```
df['Yearly equip failure'].describe()
count
        9991.000000
           0.397157
mean
std
           0.633253
           0.000000
min
25%
           0.000000
50%
           0.000000
75%
           1.000000
           4.000000
max
```

• <u>Techie</u>. Out of 9991 customers, only 16.8% of customers described themselves as technically inclined.

```
df['Techie'].value_counts()

Techie
0 8314
1 1677
Name: count, dtype: int64
```

• Port modem. Out of 9991 customers, 51.6% has a portable modem.

```
df['Port_modem'].value_counts()

Port_modem
0 5158
1 4833
Name: count, dtype: int64
```

• <u>Tablet</u>. Out of 9991 customers, only 29.9% reported owning a tablet.

```
df['Tablet'].value_counts()

Tablet
0 7006
1 2985
Name: count, dtype: int64
```

• <u>InternetService</u>. Customer's internet service provider could be fiber optics, DSL, or none. 2128 out of 9991 customers (21.3%) have neither fiber optics nor DSL. More customers have fiber optics than DSL (44.1% vs. 34.6%).

```
df['InternetService_Fiber Optic'].value_counts()

InternetService_Fiber Optic
0    5587
1    4404
Name: count, dtype: int64

df['InternetService_DSL'].value_counts()

InternetService_DSL
0    6532
1    3459
Name: count, dtype: int64
```

• Phone. 90.7% of customers own a phone.

```
df['Phone'].value_counts()

Phone
1 9058
0 933
Name: count, dtype: int64
```

• Multiple. Less than half of customers have multiple phone lines, 46.1%.

```
df['Multiple'].value_counts()

Multiple
0 5387
1 4604
Name: count, dtype: int64
```

OnlineSecurity. Just over a third of customers (35.8%) have an online security addon.

```
df['OnlineSecurity'].value_counts()

OnlineSecurity
0 6418
1 3573
Name: count, dtype: int64
```

OnlineBackup. Just under half of all customers, 45.1%, have an addon for online backup.

```
df['OnlineBackup'].value_counts()

OnlineBackup
0 5489
1 4502
Name: count, dtype: int64
```

• StreamingTV. About half of all customers (49.3%) stream TV.

```
df['StreamingTV'].value_counts()

StreamingTV
0 5068
1 4923
Name: count, dtype: int64
```

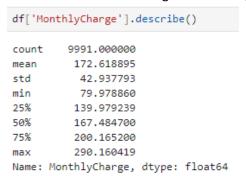
• StreamingMovies. About half of all customers (48.9%) stream movies.

```
df['StreamingMovies'].value_counts()
StreamingMovies
0 5104
1 4887
```

• <u>Tenure</u>. The average tenure of customers is 34 months. The middle 50% of customers have been with the provider between 8 and 61 months. The most tenured customer has been with the provider for 72 months, while the least tenured is 1 month.

```
df['Tenure'].describe()
count 9991.000000
         34.508248
mean
          26.442933
std
           1.000259
min
25%
          7.916107
50%
          33.196120
          61.473295
75%
          71.999280
max
Name: Tenure, dtype: float64
```

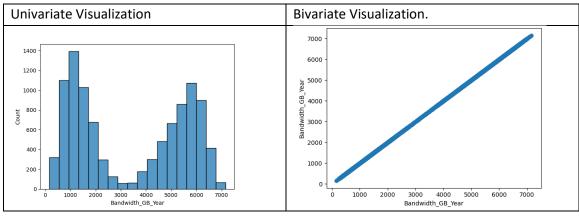
• MonthlyCharge. The average monthly charge per customer is \$172.62 (median \$167.48). The middle 50% of customers are charged between \$139.98 and \$200.17 monthly.



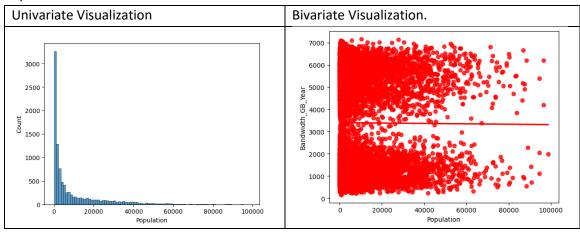
C3

Univariate and bivariate visualizations are for all variables are presented here. Bandwidth_GB_Year is the only dependent variable. Bivariate visualizations are plotted against Bandwidth_GB_Year

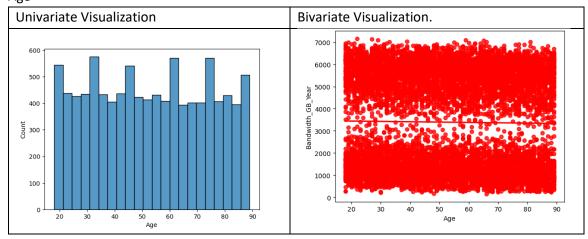
• Bandwidth_GB_Year (dependent variable).



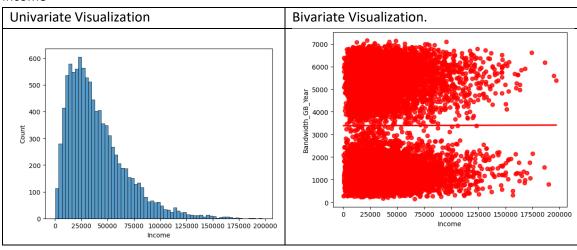
Population



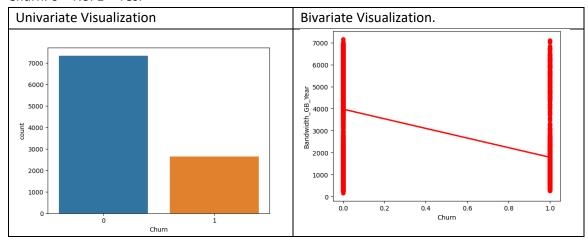
Age



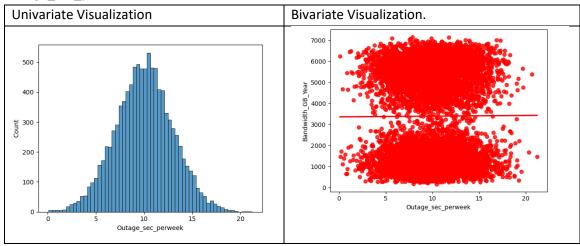
Income



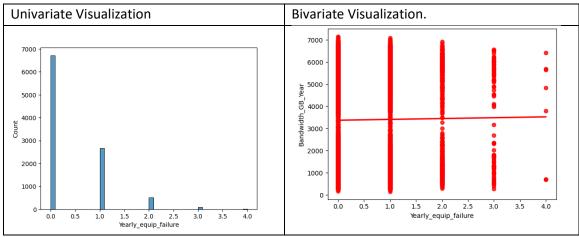
• Churn. 0 = No. 1 = Yes.



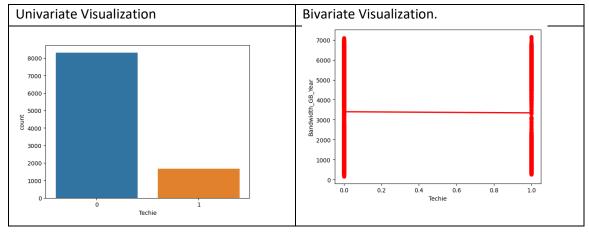
Outage_sec_perweek



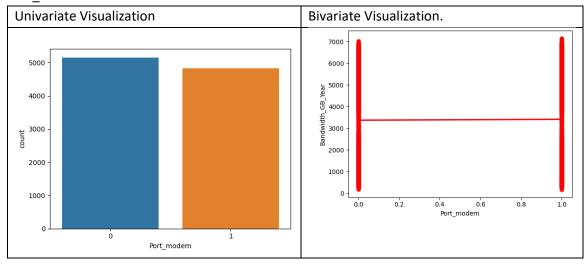
Yearly_equip_failure



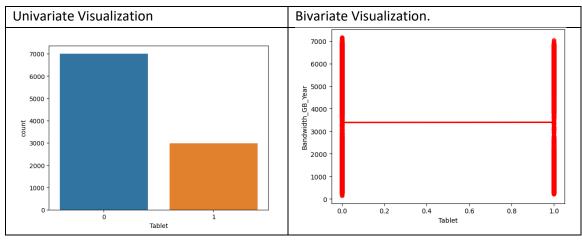
Techie. 0 = No. 1 = Yes.



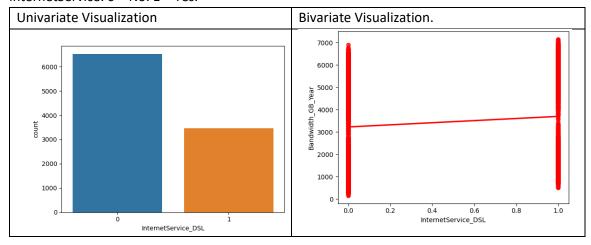
Port_modem. 0 = No. 1 = Yes.

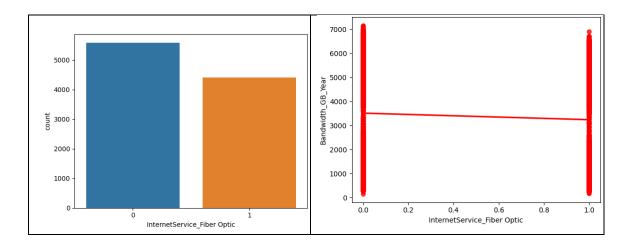


• Tablet. 0 = No. 1 = Yes.

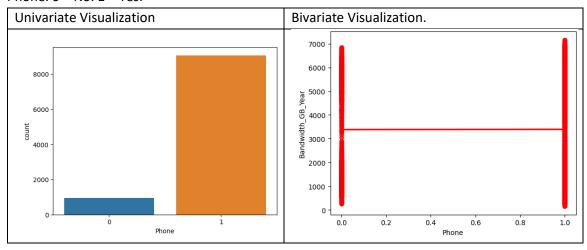


• InternetService. 0 = No. 1 = Yes.

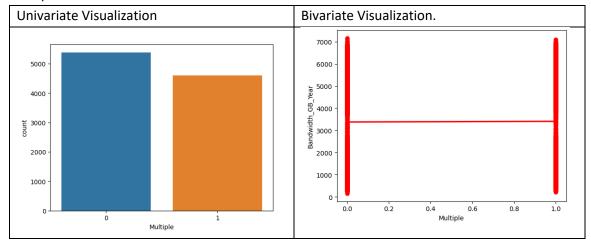




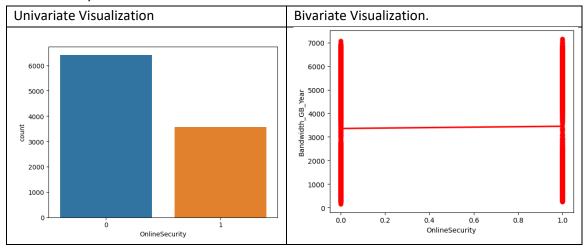
• Phone. 0 = No. 1 = Yes.



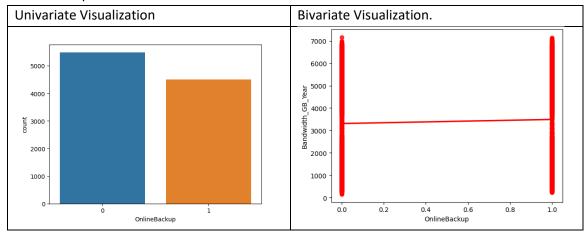
• Multiple. 0 = No. 1 = Yes.



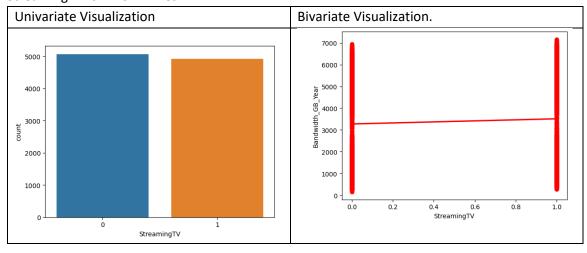
• OnlineSecurity. 0 = No. 1 = Yes.



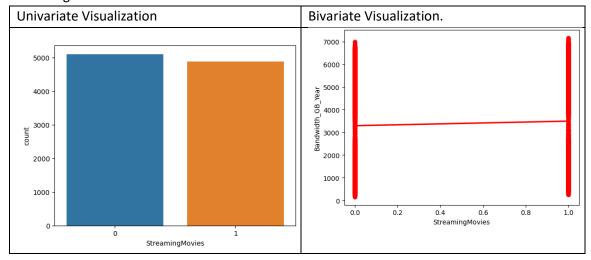
• OnlineBackup. 0 = No. 1 = Yes.



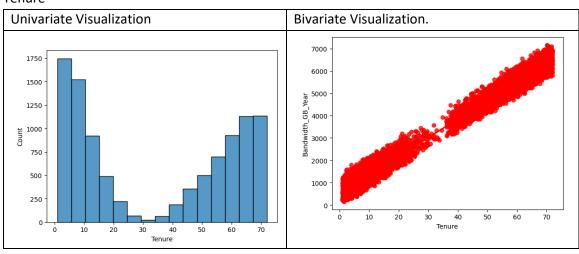
• StreamingTV. 0 = No. 1 = Yes.



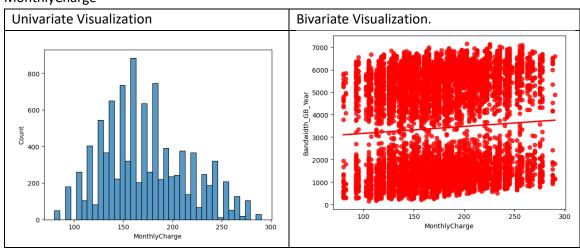
• StreamingMovies. 0 = No. 1 = Yes.



• Tenure

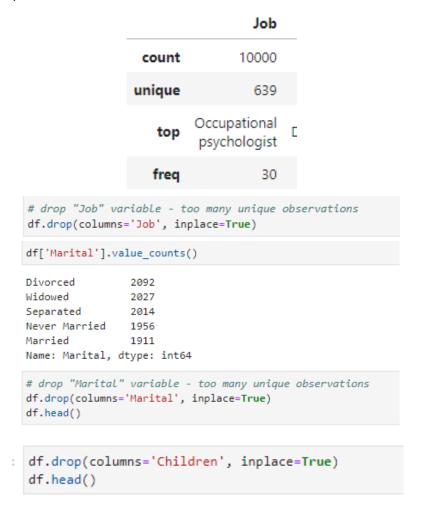


MonthlyCharge



C4

My data transformation goals included determining which variables to keep. I initially kept all variables that I thought would influence the response variable "Bandwidth_GB_Year". "Job" was one of the variables I thought to include, but there 639 unique values, so I excluded it from the analysis. I also dropped "Martial" for the same reason, too many unique values. I also dropped "Children" so the focus is just the customer, not the household.



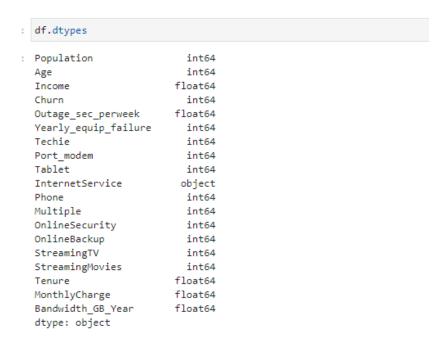
The explanatory variables used in my initial regression are:

- Population
- Age
- Income
- Churn
- Outage_sec_perweek
- Yearly_equip_failure

- Techie
- Port modem
- Tablet
- InternetService
- Phone
- Multiple
- OnlineSecurity
- OnlineBackup
- StreamingTV
- StreamingMovies
- Tenure
- MonthlyCharge

Other data transformation goals are to re-express categorical variables for use in the regression models. First I re-expressed <u>ordinal</u> categorical data types using ordinal encoding. These included variables with a "yes" or "no" value. All "yes" values were replaced with "1" and all "no" values were replaced with "0". Then I made sure the datatype changed from object to int64.

```
df['Churn'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['Churn'].value_counts()
: 0
       7341
  1
       2650
  Name: Churn, dtype: int64
 df['Techie'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['Port_modem'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
 df['Tablet'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['Phone'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['Multiple'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['OnlineSecurity'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['OnlineBackup'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['StreamingTV'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['StreamingMovies'].replace({'No' : 0, 'Yes' : 1}, inplace=True)
  df['StreamingMovies'].value counts()
  0
       5104
       4887
  Name: StreamingMovies, dtype: int64
```



Next, I re-expressed <u>nominal</u> categorical data types for the variable "InternetService" using pd.getdummies(). "InternetService" has three values, with no inherent rank.



C5

A copy of the prepared data set is submitted as a CSV file titled 'prepared_dataset_churn_D208.csv'.

D1

The initial multiple linear regression model from all independent variables that were identified in part C2 is expressed below.

```
Bandwidth\_GB\_Year = 13.59250
      +-0.00003875838 * Population
      + -3.359748 * Age
      + 0.0000002482354 * Income
      + 1.678670 * Churn
      + 0.375506 * Outage_sec_perweek
      + -0.4325477 * Yearly_equip_failure
      + - 1.330278 * Techie
      + 1.828037 * Port_modem
      + -0.4262595 * Tablet
      + 01.906635 * Phone
      + -0.3745406 * Multiple
      + 69.28261 * OnlineSecurity
      + 20.06130 * OnlineBackup
      + 89.90295 * StreamingTV
      +41.04283 * Streaming Movies
      + 41.04283 * Tenure
      +3.272196*MonthlyCharge
      + 370.03850 * InternetService_DSL
      + - 107.8163 * InternetService_FiberOptic
```

D2

I reduced the initial model, represented in part D1, to better align with the research question by engaging in the following tasks.

First, I checked for multicollinearity among the variables. This is not a feature selection procedure, but instead is a method for reducing multicollinearity, an assumption of multiple linear regression. One variable, "MonthlyCharge", had a VIF greater than 10. This variable was removed due to a high VIF factor.

Next, I ran the model without "MonthlyCharge". On this step, I engaged in a statistically based feature selection procedure by selecting variables using an iterative backwards elimination method until all variables are statistically significant with a p-value < 0.05. I removed the variable with the highest p-value first,"Yearly_equip_failure", with p=0.9820. Then I ran the model again. "Techie" now had the highest p-value of p=0.767 so it was removed. After running the model again, "Port_modem" had the largest p-value of p=0.537.

Next, all variables in the model are statistically significant with p < 0.05. The reduced model is provided in part D3.

D3

As a result of the feature selection process described in part D2, the reduced linear regression model is as follows:

```
Bandwidth\_GB\_Year = 278.3015
         + - 0.00005178 * Population
         + - 3.3205 * Age
         0.00002559 * Income
         + 15.9373 * Churn
         + 0.4132 * Outage_sec_perweek
         + - 1.6094 * Tablet
         + -5.6398 * Phone
         + 67.3540 * Multiple
         + 79.2940 * OnlineSecurity
         + 93.3939 * OnlineBackup
         + 225.0095 * StreamingTV
         + 208.8257 * StreamingMovies
         +81.9902 * Tenure
         +411.1733 * InternetService_DSL
         + -1.6165 * InternetService\_FiberOptic
```

• OLS Results for the <u>Initial Linear Regression Model</u>:

OLS Regression Results

	OLS	OLS Regression Results					
Dep. Variable	s Bandwid	dth_GB_Year	R-	squared:	0.99	99	
Mode	l:	OLS	Adj. R-	squared:	0.99	99	
Method	l: Le	ast Squares	F-	statistic:	3.906e+0	05	
Date	Sun, 1	9 Nov 2023	Prob (F-s	tatistic):	0.0	00	
Time	8	17:12:39	Log-Lik	elihood:	-5796	2.	
No. Observations	:	9991		AIC:	1.160e+0	05	
Df Residuals	:	9971		BIC:	1.161e+0	05	
Df Mode	l:	19					
Covariance Type	*	nonrobust					
		coef	std err		t P> t	[0.025	0.975]
	Intercept	13.5925	7.451	1.82	4 0.068	-1.012	28.197
Monti	nlyCharge	3.2722	0.066	49.95	5 0.000	3.144	3.401
P	opulation	-3.876e-05	5.59e-05	-0.69	3 0.488	-0.000	7.09e-05
	Age	-3.3597	0.039	-86.69	3 0.000	-3.436	-3.284
	Income	2.482e-07	2.88e-05	0.00	9 0.993	-5.62e-05	5.67e-05
	Churn	1.6787	2.377	0.70	5 0.480	-2.980	6.337
Outage_sec	perweek	0.3755	0.269	1.39	4 0.163	-0.153	0.904
Yearly_equ	ip_failure	0.4325	1.267	0.34	2 0.733	-2.050	2.915
	Techie	-1.3303	2.153	-0.61	8 0.537	-5.551	2.891
Por	t_modem	1.8280	1.605	1.13	9 0.255	-1.318	4.974
	Tablet	-0.4263	1.753	-0.24	3 0.808	-3.863	3.011
	Phone	-1.9066	2.759	-0.69	1 0.489	-7.314	3.501
	Multiple	-37.4541	2.657	-14.09	7 0.000	-42.662	-32.246
Onlin	eSecurity	69.2826	1.686	41.09	5 0.000	65.978	72.587
Onli	neBackup	20.0613	2.184	9.18	6 0.000	15.780	24.342
Stre	eamingTV	89.9030	3.182	28.25	0.000	83.665	96.141
Streami	ngMovies	41.0428	3.771	10.88	2 0.000	33.650	48.436
	Tenure	81.9007	0.036	2280.23	7 0.000	81.830	81.971
InternetSe	rvice_DSL	370.3950	2.365	156.62	8 0.000	365.759	375.030
InternetService_F	iberOptic	-107.8163	3.000	-35.94	2 0.000	-113.696	-101.936
Omnibus:	1037.448	Durbin-W	atson:	2.005			
Prob(Omnibus):	0.000	Jarque-Ber	a (JB): 1	455.422			
Skew:	0.820	Pro	b(JB):	0.00			

Cond. No. 5.27e+05

Kurtosis:

3.897

• OLS Results <u>Reduced Linear Regression Model</u>:

OLS Regression Results

Dep. Variable:	Bandwidth_GB_Year	R-squared:	0.998
Model:	OLS	Adj. R-squared:	0.998
Method:	Least Squares	F-statistic:	3.957e+05
Date:	Sun, 19 Nov 2023	Prob (F-statistic):	0.00
Time:	17:18:42	Log-Likelihood:	-59079.
No. Observations:	9991	AIC:	1.182e+05
Df Residuals:	9975	BIC:	1.183e+05
Df Model:	15		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	278.3015	5.771	48.226	0.000	266.990	289.613
Population	-5.178e-05	6.25e-05	-0.828	0.407	-0.000	7.08e-05
Age	-3.3205	0.043	-76.661	0.000	-3.405	-3.236
Income	2.559e-05	3.22e-05	0.795	0.426	-3.75e-05	8.87e-05
Churn	15.9373	2.629	6.062	0.000	10.784	21.091
Outage_sec_perweek	0.4132	0.301	1.372	0.170	-0.177	1.004
Tablet	-1.6094	1.960	-0.821	0.412	-5.451	2.232
Phone	-5.6398	3.083	-1.829	0.067	-11.683	0.403
Multiple	67.3540	1.822	36.971	0.000	63.783	70.925
OnlineSecurity	79.2940	1.871	42.380	0.000	75.626	82.962
OnlineBackup	93.3939	1.808	51.667	0.000	89.851	96.937
StreamingTV	225.0095	1.874	120.068	0.000	221.336	228.683
StreamingMovies	208.8257	1.916	108.979	0.000	205.070	212.582
Tenure	81.9902	0.040	2045.848	0.000	81.912	82.069
InternetService_DSL	411.1733	2.481	165.717	0.000	406.310	416.037
InternetService_FiberOptic	-1.6165	2.366	-0.683	0.495	-6.254	3.022

Omnibus:	606.095	Durbin-Watson:	1.982
Prob(Omnibus):	0.000	Jarque-Bera (JB):	729.668
Skew:	0.619	Prob(JB):	3.59e-159
Kurtosis:	3.468	Cond. No.	3.36e+05

E1

In my data analysis process I developed a multiple linear regression model to determine what factors influence how much bandwidth customers use. I will use the adjusted R^2 value as a model evaluation metric to make a comparison of the initial and reduced multiple linear regression model. For the initial model adjusted $R^2=0.999$ and for the reduced model adjusted $R^2=0.998$. Both models explain over 99% of the variation, while taking into account the number of predictor variables (Van den Broeck, nd).

Initial Model Calculations

R-squared:	0.999
Adj. R-squared:	0.999
F-statistic:	3.906e+05
Prob (F-statistic):	0.00
Log-Likelihood:	-57962.
AIC:	1.160e+05
BIC:	1.161e+05

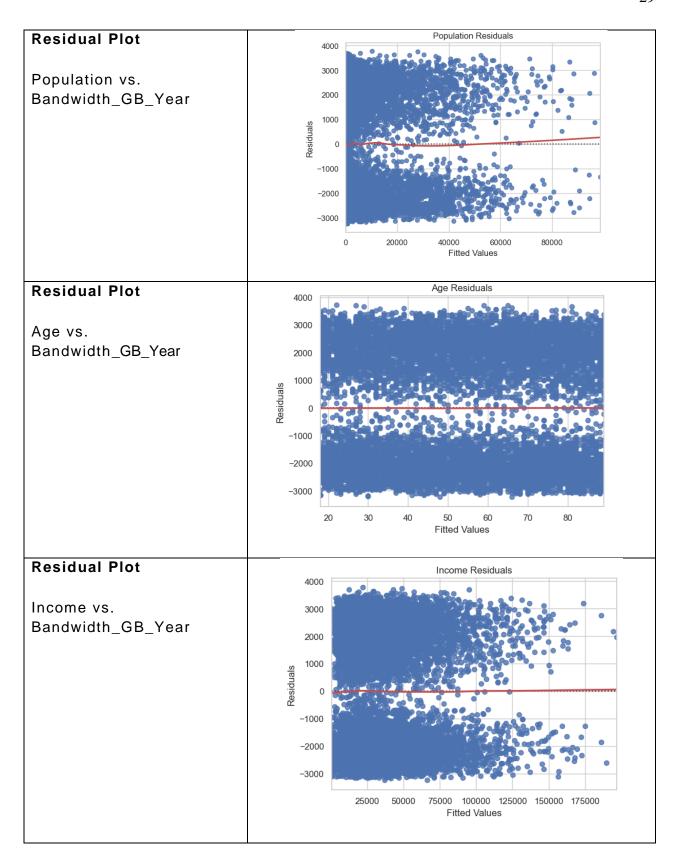
Reduced Model Calculations

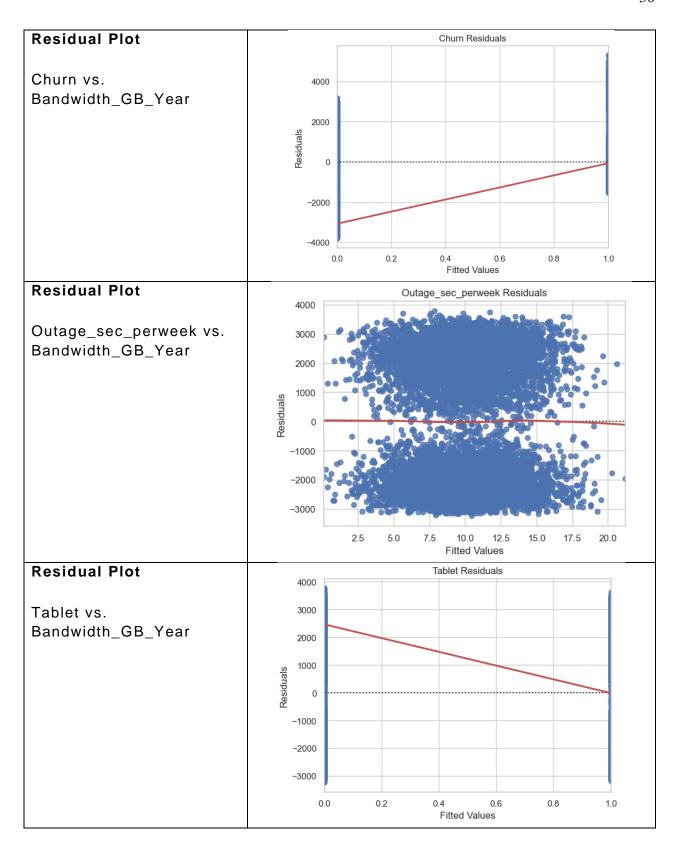
R-squared:	0.998
Adj. R-squared:	0.998
F-statistic:	3.957e+05
Prob (F-statistic):	0.00
Log-Likelihood:	-59079.
AIC:	1.182e+05
BIC:	1.183e+05

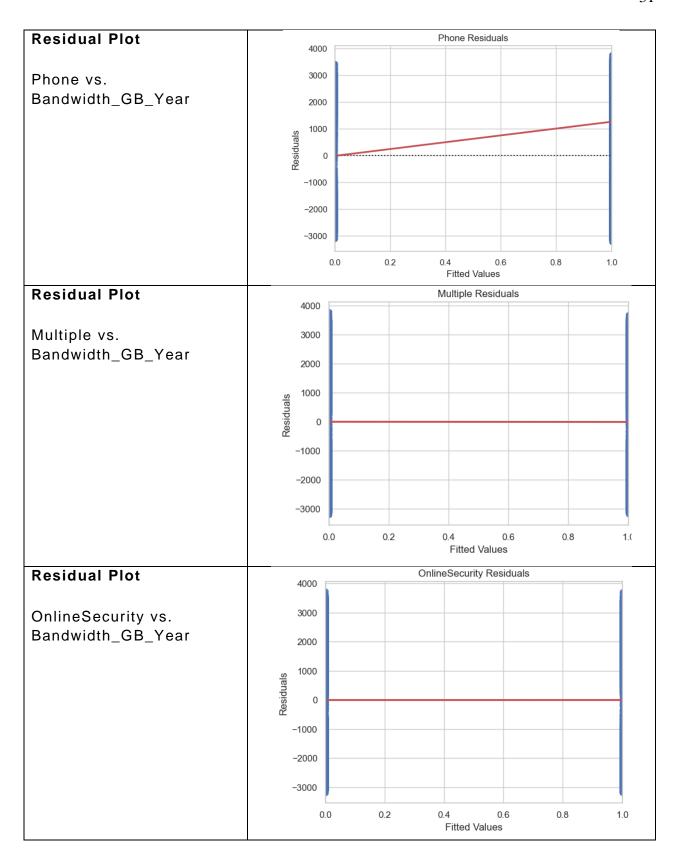
E2

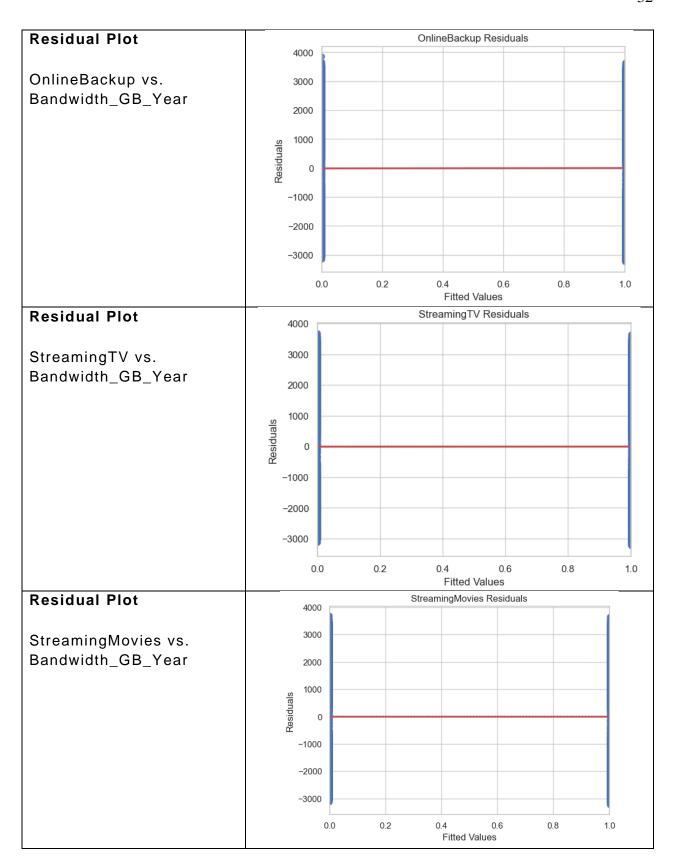
The output and all calculations of the analysis I performed for the reduced model are outlined below as well as a residual plot for each independent variable vs. Bandwidth_GB_Year

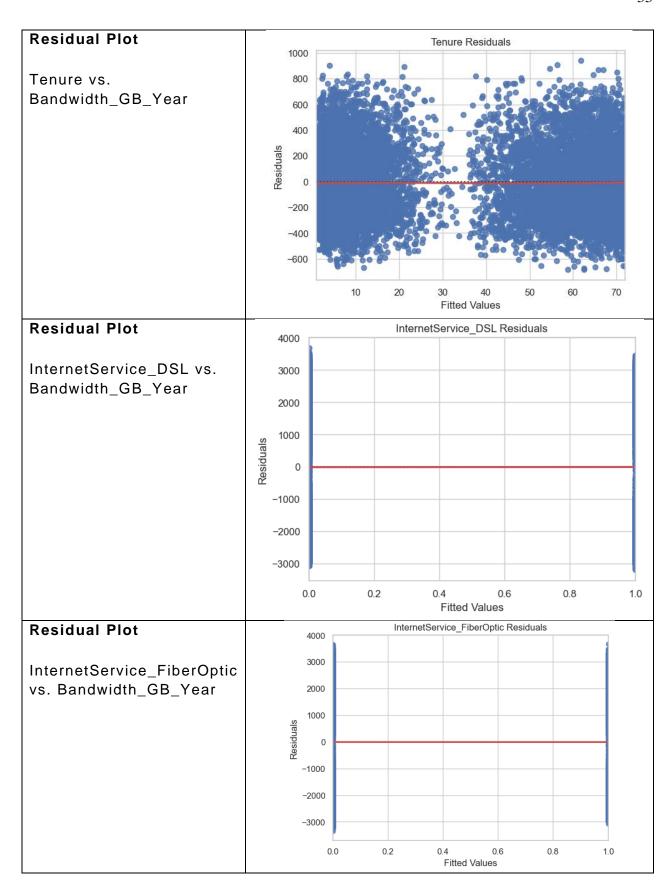
Reduced Model		R-squared:	0.998
Calculations		Adj. R-squared:	0.998
	men - model neduced4 men perid	F-statistic:	3.957e+05
	<pre>mse = model_reduced4.mse_resid print('mse: ' , mse)</pre>	Prob (F-statistic):	0.00
	rse = np.sqrt(mse) print('rse: ', rse) mse: 8023.015882801472 rse: 89.57128938896365	Log-Likelihood:	-59079.
		AIC:	1.182e+05
		BIC:	1.183e+05











E3

A copy of the code used to support the implementation of the linear regression model is submitted. There are two code files, "D208 PA Task 1 – Code File Part 1" and "D208 PA Task 1 – Code File Part 3". The Part 3 file uses the cleaned data set created from Part 1 to perform multiple linear regression.

F1

Here I provide a summary of my findings and assumptions by discussing the regression equation for the reduced model, an interpretation of the coefficients of the reduced model, the statistical and practical significance of the recued model, and the limitations of the data analysis.

The regression equation is given by:

```
Bandwidth\_GB\_Year = 278.3015
      -0.00005178 * Population
      + - 3.3205 * Age
      0.00002559 * Income
      + 15.9373 * Churn
      + 0.4132 * Outage_sec_perweek
      + - 1.6094 * Tablet
      + - 5.6398 * Phone
      + 67.3540 * Multiple
      + 79.2940 * OnlineSecurity
      + 93.3939 * OnlineBackup
      + 225.0095 * StreamingTV
      +208.8257 * Streaming Movies
      +81.9902 * Tenure
      + 411.1733 * InternetService_DSL
      + - 1.6165 * InternetService_FiberOptic
```

• An interpretation for the coefficients of the reduced model is given below:

Coefficient * Variable	Interpretation
Intercept = 278.3015	The y-intercept represents when all other
	independent variables are 0, the amount of
	bandwidth used per year on average is 278.3015.
-0.00005178*Population	For every one person increase in population
	within a mile radius of the customer, the
	customer uses about 0 more GB in bandwidth per
	year on average.

2 2 2 2 2	
-3.3205 * Age	For every one year increase in age, the customer uses about 3.3 less GB in bandwidth per year on average.
0.00002559 * Income	For every one dollar increase in income, the customer uses about 0 more GB in bandwidth per year on average.
15.9373 * Churn	Customers who churn use about 16 more GB in bandwidth per year on average than customers who don't.
0.4132 * Outage_sec_perweek	For every one second increase in average weekly outages, customers use about 0.4 more GB per year on average.
-1.6094 * Tablet	Customers who have a tablet use about 1.6 less GB in bandwidth per year on average than customers who don't.
-5.6398 * Phone	Customers who have a phone service use about 5.6 less GB in bandwidth per year on average than customers who don't.
67.3540 * Multiple	Customers who have multiple lines use about 67.4 more GB in bandwidth per year on average than customers who don't.
79.2940 * OnlineSecurity	Customers who have an online security add-on use about 79.3 more GB in bandwidth per year on average than customers who don't.
93.3939 * OnlineBackup	Customers who have an online backup add-on use about 93.4 more GB in bandwidth per year on average than customers who don't.
225.0095 * StreamingTV	Customers who have streaming TV use about 225 more GB in bandwidth per year on average than customers who don't.
208.8257 * StreamingMovies	Customers who have streaming movies use about 208.8 more GB in bandwidth per year on average than customers who don't.
81.9902 * Tenure	For every one month increase in tenure, customers use about 82 more GB in bandwidthper year on average.
411.1733 * InternetService_DSL	Customers who have DSL internet use about 411.2 more GB in bandwidth per year on average than customers who don't.
-1.6165 * InternetService_FiberOptic	Customers who have Fiber Optic internet use about 1.6 less GB in bandwidth per year on average than customers who don't.

The statistical and practical significance of the reduced model are discussed here.

Each coefficient in the reduced model is statistically significant because the p-value is less than 0.05 for each. RSE is 89.6, meaning the difference between predicted values and observed values is typically 89.6.

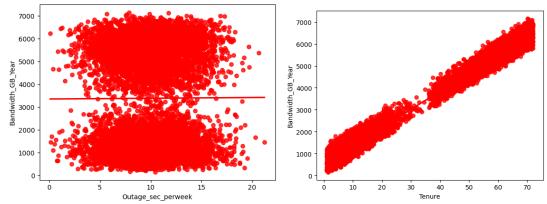
The practical significance of the model could be lacking due to the higher number of independent variables used—there are fifteen independent variables. This can make the model hard to interpret and potentially less practical to use in the real world. Also, population and income have a coefficient essentially equal to 0, so these factors have hardly any influence on how much bandwidth a customer is using. Also, "Outage_sec_perweek" might not be a very useful variable for practical applications of the model.

• The limitations in the analysis are as follows (Middleton, nd):

First, not all assumptions of linear regression are met.

I assume that the observations in this data set are independent and random.

Multiple linear regression assumes that each independent variable is linearly related to the dependent variable. In part C3, a scatterplot was made to visualize if such a linear relationship existed. Some data when paired with "Bandwidth_GB_Year", such as "Outage_sec_perweek", was not linear, especially when compared with "Tenure".



Independent variables should not be highly correlated with one another. As described in part D2, a VIF was calculated for each variable to check for multicollinearity and only one factor was removed as a result.

Normality of residuals is an additional assumption of linear regression. The Prob(Omnibus) Test yields a value of 0, and since it is smaller than 0.05, this indicates that the residuals are not

normally distributed. Furthermore, Prob(Jarque-Bera) Test also concludes that the residuals are not normally distributed because it also yielded a value less 0.05 ("Assumptions of Multiple Linear Regression", 2023).

1.982	Durbin-Watson:	606.095	Omnibus:
729.668	Jarque-Bera (JB):	0.000	Prob(Omnibus):
3.59e-159	Prob(JB):	0.619	Skew:
3.36e+05	Cond. No.	3.468	Kurtosis:

The last assumption is homoscedasticity. The residual plots, found in part E2, indicate that data are homoscedastic.

Additionally, the data in the reduced model still contains outliers, which affects the model. The model has a very high \mathbb{R}^2 value, which means this model is overfitting the data.

F₂

My recommended course of action based on the results my analysis is as follows:

Continue to refine the model to make it more practically significant by reducing the number of independent variables, using interactions between variables, transforming variables to make data more linear, and removing more outliers.

However, the reduced linear regression model does offer some insight into which factors are influencing customers' use of bandwidth. The aim of the research question was to help the company optimize the services they offer to customers based on factors that influence bandwidth. We can see that population and income have very little influence on the amount of bandwidth used due to their near 0 coefficients. Also, customers who have a DSL internet service use quite a lot more data on average than customers who don't.

Additionally, the company may consider selling a data plan add-on for their services with different thresholds of data. For example, if a customer streams TV or movies, they are using 209-225 more GB per year on average than customers who don't. The company also sells a data package to the customer that is equipped to handle the streaming services.

G

A Panopto video recording is provided in the submission of this performance assessment: https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=eb06ab32-ff8e-4e1f-b834-b0c00023105f

Н

Web sources used to acquire segments of code to support the application:

"Data Science – Statistics Correlation Matrix." W3 Schools.

www.w3schools.com/datascience/ds stat correlation matrix.asp

"Detecting Multicollinearity with VIF – Python." Geeks for Geeks. <u>www.geeksforgeeks.org/detecting-</u> multicollinearity-with-vif-python/

"seaborn.residplot." seaborn. seaborn.residplot. seaborn.pydata.org/generated/seaborn.residplot.html

Van den Broeck, Maarten. "Introduction to Regression with statsmodels in Python." Datacamp, app.datacamp.com/learn/courses/introduction-to-regression-with-statsmodels-in-python

Van den Broeck, Maarten. "Intermediate Regression with statsmodels in Python." Datacamp, app.datacamp.com/learn/courses/intermediate-regression-with-statsmodels-in-python



I acknowledged sources, using in-text citations and references, for content that is quoted, paraphrased, or summarized.

References:

"Assumptions of Multiple Linear Regression". Complete Dissertation by Statistics Solutions. 2023.

 $\underline{www.statistics solutions.com/free-resources/directory-of-statistical-analyses/assumptions-of-multiple-linear-}$

regression/#:~:text=The%20last%20assumption%20of%20multiple,)%2C%20the%20data%20is%20heteroscedastic.

Middleton, Keiona. "Getting Started with D208 Part I." Predictive Modeling – D208, nd, westerngovernorsuniversity.sharepoint.com/sites/DataScienceTeam/Shared%20Documents/For

ms/AllItems.aspx?csf=1&web=1&e=9ccodm&cid=3911c4ec%2D7d83%2D4ba9%2Da54b%2D7da
9cbb5d38f&FolderCTID=0x01200022092E63FD85A64A8ABFB4F5AEA4839A&id=%2Fsites%2FDat
aScienceTeam%2FShared%20Documents%2FGraduate%20Team%2FD208%2FStudent%20Facing
%20Resources%2FDr%2E%20Middleton%20Getting%20Started%20with%20D208%28Part%20I%
29COIT%2Epdf&parent=%2Fsites%2FDataScienceTeam%2FShared%20Documents%2FGraduate
%20Team%2FD208%2FStudent%20Facing%20Resources

Middleton, Keiona. "Getting Started with D208 Part II." Predictive Modeling - D208, nd,

westerngovernorsuniversity.sharepoint.com/sites/DataScienceTeam/Shared%20Documents/Forms/AllItems.aspx?csf=1&web=1&e=9ccodm&cid=3911c4ec%2D7d83%2D4ba9%2Da54b%2D7da9cbb5d38f&FolderCTID=0x01200022092E63FD85A64A8ABFB4F5AEA4839A&id=%2Fsites%2FDataScienceTeam%2FShared%20Documents%2FGraduate%20Team%2FD208%2FStudent%20Facing%20Resources%2FDr%2E%20Middleton%20Getting%20Started%20with%20D208%20%28Part%20II%29COIT%2Epdf&parent=%2Fsites%2FDataScienceTeam%2FShared%20Documents%2FGraduate%20Team%2FD208%2FStudent%20Facing%20Resources

Nair, Aashish. "Targeting Multicollinearity With Python." Medium. December 6, 2021, towardsdatascience.com/targeting-multicollinearity-with-python-3bd3b4088d0b

Van den Broeck, Maarten. "Introduction to Regression with statsmodels in Python." Datacamp, app.datacamp.com/learn/courses/introduction-to-regression-with-statsmodels-in-python

Van den Broeck, Maarten. "Intermediate Regression with statsmodels in Python." Datacamp, app.datacamp.com/learn/courses/intermediate-regression-with-statsmodels-in-python



Professional communication is demonstrated in the content and presentation of my Performance Assessment.